

WID3014: Practical AI

Exercises based on Lectures for Week 10

(Deadline 23rd December 2024, 1500hr)

1 Distributed Gradient Descent in Python: A Tutorial

Introduction

Gradient descent is a widely used optimization algorithm for training machine learning models. In this tutorial, we will learn how to implement a **distributed gradient descent** (GD) algorithm for training a simple **linear regression** model. The goal of distributed gradient descent is to split the work of calculating gradients across multiple processes or machines to accelerate model training.

For this tutorial, we will use Python's multiprocessing library to simulate distributed gradient descent on a single machine (i.e., a multi-core CPU). This approach can also be extended to more complex distributed systems, but for simplicity, we will simulate the process locally.

Prerequisites

1. **Google Colab** or a **local machine** with Python and numpy installed.
2. **Understanding of gradient descent** and **linear regression**.

Step 1: Setup and Sample Data

Let's start by creating a synthetic dataset that we will use to train a simple linear regression model. Our model will have the equation:

$$y = w \cdot x + b$$

Where:

- **w** is the weight (slope),
- **b** is the bias (intercept).

We'll generate synthetic data with some noise to make it more realistic.

```
import numpy as np
import matplotlib.pyplot as plt

# Create synthetic data
np.random.seed(42)
X = np.random.rand(1000, 1) # 1000 samples, 1 feature
y = 2 * X + np.random.randn(1000, 1) * 0.1 # y = 2*x + noise

# Visualize the data
plt.scatter(X, y, c='blue', label='Data points')
```

```
plt.xlabel('X')
plt.ylabel('y')
plt.title('Sample Dataset for Linear Regression')
plt.show()
```

Step 2: Gradient Descent Function

We now implement the gradient descent algorithm. The goal is to minimize the Mean Squared Error (MSE) between the predicted values and the actual values. For linear regression, the loss function (MSE) is:

$$MSE = \frac{1}{m} \sum_{i=1}^m (y_i - \hat{y}_i)^2$$

Where:

- m is the number of samples
- \hat{y}_i is the predicted value

The gradient of the MSE with respect to w and b is:

$$\frac{\partial MSE}{\partial w} = \frac{2}{m} \sum_{i=1}^m (y_i - \hat{y}_i) \cdot x_i$$

$$\frac{\partial MSE}{\partial b} = \frac{2}{m} \sum_{i=1}^m (y_i - \hat{y}_i)$$

The updates for the parameters are:

$$w = w - \eta \cdot \frac{\partial MSE}{\partial w}$$

$$b = b - \eta \cdot \frac{\partial MSE}{\partial b}$$

Where η is the learning rate:

```
def compute_mse(X, y, w, b):
    m = len(y)
    predictions = X.dot(w) + b
    mse = (1/m) * np.sum((predictions - y) ** 2)
    return mse

def compute_gradients(X, y, w, b):
```

```

m = len(y)
predictions = X.dot(w) + b
dw = (2/m) * X.T.dot(predictions - y)
db = (2/m) * np.sum(predictions - y)
return dw, db

```

Step 3: Simulate Distributed Gradient Descent

We will now implement the core of distributed gradient descent using Python's multiprocessing library. The dataset will be divided into smaller chunks, and each process will compute the gradients on its portion of the data. The results will then be aggregated to update the parameters.

Here is the implementation:

```

import multiprocessing as mp

# Function to be run by each process
def gradient_descent_process(rank, X_split, y_split, w, b, result_queue):
    dw, db = compute_gradients(X_split, y_split, w, b) # Compute gradients on local data
    result_queue.put((dw, db)) # Place gradients in result queue

# Function to run distributed gradient descent
def distributed_gradient_descent(X, y, learning_rate=0.01, epochs=100, num_processes=4):
    # Split the data across different processes
    data_split = np.array_split(X, num_processes)
    target_split = np.array_split(y, num_processes)

    # Initialize parameters (weights and bias)
    w = np.zeros((X.shape[1], 1)) # Initialize weight to zero
    b = np.zeros((1, 1)) # Initialize bias to zero

    # Queue for collecting results from processes
    result_queue = mp.Queue()

    # Run gradient descent
    for epoch in range(epochs):
        processes = []

        # Start multiple processes for gradient computation
        for rank in range(num_processes):
            p = mp.Process(target=gradient_descent_process, args=(rank, data_split[rank],
            target_split[rank], w, b, result_queue))
            processes.append(p)
            p.start()

        # Collect results from all processes
        dw_total = np.zeros_like(w)
        db_total = np.zeros_like(b)
        for p in processes:

```

```

p.join() # Wait for all processes to finish
while not result_queue.empty():
    dw, db = result_queue.get()
    dw_total += dw
    db_total += db

# Update weights and bias
w -= learning_rate * dw_total
b -= learning_rate * db_total

# Print the progress
if epoch % 10 == 0:
    mse = compute_mse(X, y, w, b)
    print(f"Epoch {epoch}: MSE = {mse}")

return w, b

```

Step 4: Running the Code

Now, let's run the distributed gradient descent algorithm and observe the results.

```

# Run distributed gradient descent
w_final, b_final = distributed_gradient_descent(X, y)

# Print the final parameters
print("Final weights:", w_final)
print("Final bias:", b_final)

```

Explanation of the Code:

1. **Data Splitting:** The dataset X and y is divided into num_processes parts, where each part is handled by a different process.
2. **Gradient Calculation:** Each process computes the gradients based on its subset of data.
3. **Result Collection:** The gradients are collected into a queue, which is then processed by the main thread to sum the gradients.
4. **Parameter Update:** The weights and bias are updated using the total gradients computed across all processes.
5. **Final Output:** After all epochs, the final learned parameters (weights and bias) are printed.

Step 5: Run the Code in Google Colab or Locally

- **Google Colab:** You can run the code directly in a Google Colab notebook. Colab supports multi-processing, and this code should run without modification.
- **Local Machine:** If you are running this on a local machine with multiple CPU cores, you can execute the code as-is, and it will take advantage of the available cores for parallel processing.

To run the code locally, save it in a .py file and execute:

```
python distributed_gradient_descent.py
```

2 Questions

1. What is the purpose of using multiple processes in this implementation?
2. Why do we use a Queue in the code?
3. What would happen if we did not aggregate the gradients across all processes?
4. What is the role of the learning_rate parameter?
5. How can the number of processes (num_processes) affect the training process?
6. How does scaling data storage relate to distributed gradient descent?
7. How would you handle very large datasets that don't fit in memory during distributed gradient descent?
8. What are some common distributed storage systems used for large-scale machine learning workloads?
9. What are the benefits of using distributed data processing frameworks like Apache Spark or Dask in gradient descent?
10. How would you distribute data storage and processing when training a model on multiple machines?
11. What is the role of batch processing in distributed gradient descent when dealing with large datasets?
12. How can you optimize the distributed gradient descent process when working with huge datasets?
13. How can data compression be used to handle large datasets in distributed gradient descent?
14. What challenges might arise when scaling data storage and processing in distributed systems?